

Numbers Claude gives about your business, with the work shown

Claude skills for the numbers a business asks every month, from Shopify, Stripe, Amazon or another export, and for checking the report someone else sends. Tested on nine synthetic datasets; six of them (stores D and E, a Stripe account, an Amazon seller account, a subscription business on Stripe Billing and an agency's report) were built by separate Claude sessions that never saw the skills, each with its own answer key.

The step that matters: the owner's definitions, written down once

The owner's month-end question on the Stripe, the Amazon and the SaaS account, August and September, three runs each, graded to the cent against the answer keys. Asked plainly, Opus answers on readings of its own (the month in local time, payouts by the day they reached the bank, the reserve inside what is owed).

	STRIPE (11 LINES)	AMAZON (10 LINES)	SAAS (12 LINES)
The question alone	3–4 of 11	6–7 of 10	2–4 of 12
The question and a definitions file	10 of 11 every run	10 of 10 every run	5–6 of 12
The question and a pinned calculation	10 of 11 every run	10 of 10 every run	7 of 12 every run*

What the pin adds over the file: the same figure in every run (on Stripe the file's tax line moved by cents), 6–16 turns against 9–42, and its own check on every answer, which passed 13 of 18 and flagged the others, mostly for leaving out the open questions about Claude's readings. The line both miss on Stripe is sales tax, by cents. On the SaaS account the file gets the money right and not the subscriptions (the export gives each one's status on the day it was taken; the pin rebuilt it for the month's end), and no run reaches MRR to the cent: the answer key knows when a seat change or a retention coupon took effect, which no export records. *1 run of six stopped and asked for the file of exchange rates the setup had written.

Where the skills did not beat plain Opus

	WITH THE SKILLS	PLAIN OPUS
Store E, built apart, the Shopify skills unchanged: answers passing the five graders any right answer passes	6/9	9/9
An agency's report with 7 planted mistakes: mistakes caught, three runs	6, 6, 6	6, 6, 6
The same report: right verdicts on its 22 figures	16, 16, 16	16, 16, 16

On store E the scripts were short on gift cards sold; on the report the skill took about 1.6 times the turns. Both are told on their pages, with every answer.

The first three stores

A strong model rarely gets the arithmetic wrong any more. What goes wrong is the context: which orders count as sales, whether revenue holds VAT, which month a refund or a payout belongs to. Three skills for a Shopify store's month-end ask the owner's definitions once, compute every figure with a script, check the answer against the script's results, name every order left out, and print under every answer how it was counted. A bench then asked the owner's questions on three synthetic stores, the third at a real store's size (4,470 orders over three months), with the skills and without them, and read every answer. A fourth store, built apart from the skills by a separate Claude session, tested them against its own truth (below).

What the owner can see (Opus 5.5, three runs per question, 27 runs per arm)

	WITH THE SKILLS	WITHOUT
Passed every grader	27/27	26/27
On the skill's stated basis	27/27	13/27
Everything the owner should see, named	27/27	21/27
The skill's "How this was counted" section	25/27	0/27
Says it added the figures up by hand	0/27	9/27

STORE	PASSED, WITH	WITHOUT	ALL NAMED, WITH	WITHOUT
Store A	9/9	9/9	9/9	9/9
Store B	9/9	9/9	9/9	9/9
Store C	9/9	8/9	9/9	3/9

"Passed" is every grader: the figures on the summary and payout questions, and on margins the products that must be named (no margin figure is graded). "How this was counted" is the skill's own heading; plain answers often explain their counting in their own words.

Without the skills on store C the model had a shell and wrote its own code. It read the owner's definitions file and applied it, and it passed every grader in 8 of 9 runs; in the ninth its own code labelled a total as after discounts when it was before them, and August revenue came out about £1,000 high. It named the orders behind its counts in none of the 6 summary and payout runs: not the 85 orders left out of August's sales, not the 400 payout items (on the margins question it named every product it had to).

The answer zoo: it ran, and the answer is wrong

24 defects planted in correct answers (a wrong figure, the wrong rows counted, a basis named wrong, a right figure under another's name, a lost caveat, a claim no figure supports, the question not answered). The number check stopped **6**, the coverage check **5**; the other 13 went to reviewer models, three runs each.

REVIEWER	NAMED THE DEFECT	OBJECTED TO AN UNTOUCHED ANSWER
Sonnet 5	13 of 13	0 of 9
Opus 5.5	13 of 13	3 of 9

The objections to untouched answers were about what the answer left unsaid or put loosely (whether unshipped orders count, how much VAT the net still holds). I agree with them; the answers tested here were not changed after this run.

Store D: a store built apart from the skills

Stores A to C were built by the same work that wrote the skills, so a wrong idea of Shopify's exports could sit in both and pass. Store D was built by a separate Claude session that never saw the skills: 30,251 orders from June to August in USD, EUR, GBP and CAD, 68 traps, and a truth worked out from its own records. Only the figures that mean the same in its truth and in the skills are judged; [SHOPIFY.md](#) lists every one.

THE SKILLS' SCRIPTS AGAINST THE STORE'S TRUTH	FIRST RUN (0.6.1)	NOW (0.10.1)
Figures that match	18 of 50	55 of 56
Figures that differ	16	1
Figures not reported	16	0
OPUS 5.5, 3 RUNS PER QUESTION: ANSWERS THAT PASSED EVERY GRADER (WIDENED)		
	WITH THE SKILLS	WITHOUT
June's bank	3/3	0/3
August's sales	3/3	0/3
July's margins	2/3	1/3
All 3 questions (graders as first written; as widened after reading)	4/9 (8)	0/9 (1)
July's margins again, the skills at 0.8	2/3 (3)	
All 3 questions, the skills at 0.10.1	7/9 (9)	

The first run was blind; the versions after it were developed against this same store. Two graders were widened after the answers were read: the report an answer saved is read with its chat (only the skill side saves one) and renamed SKUs may be named as a family. The graders look for the figures the skills give, so they favour the skill side. In two payout answers at 0.10.1 and one margins answer at 0.8.0 the plugin's own answer check did not pass the final text (figures typed by hand); SHOPIFY.md keeps its warning with each.

June: all 6 runs found the \$572,456.23 that reached the bank. With the skills, 3 of 3 also gave the \$105,011.68 of May's sales in June's first payouts and 3 of 3 the \$117,366.90 still in transit at the month's end. Without them, 1 of 3 gave the \$105,011.68 and 0 of 3 the \$117,366.90, though all 3 said in words that May's money came in June's first payouts (read by hand; see SHOPIFY.md). The skills' own first run missed the same money: its scripts put June's bank at \$467,444.55, \$105,011.68 short. The one run with the skills that did not pass, on July's margins, stopped to ask for a VAT rate; 0.8 takes each order's own tax instead.

The skill zoo: does the right skill run

11 defects in the summary skill's frontmatter or files, each in its own copy of the plugin, then `claude plugin validate --strict`, the linter and a routing eval (six questions, three runs each). The last column: how often the summary skill ran on its own question, asked plainly and indirectly.

DEFECT	VALIDATE	LINTER	SUMMARY SKILL RAN
none	passed	silent	6/6
vague_description	passed	W01, W02, W05	6/6
no_trigger	passed	W01, W05	6/6
overlaps_margins	passed	W04, W05	6/6
too_broad	passed	W01, W05	6/6
name_not_kebab	passed	E04	6/6

name_differs_from_folder	passed	E04	6/6
description_field_typo	flagged	E03, E05	0/6
frontmatter_not_first	flagged	E01	0/6
model_invocation_disabled	passed	W06	0/6
trigger_past_the_limit	passed	E06	6/6
script_missing	passed	E07	6/6

Checking the checks

After each run every failure was read, and so were the passes, except the checks a reader judges at 0.10.1. Each finding changed the tools:

- **A definition the skill, its truth and its graders shared:** a model without the skill reconciled to the bank; it was right.
- **A question that said less than the script did:** "test orders" meant a tag to the owner, a gateway too to the script.
- **Reviewers right about the reference answers:** most objections were a definition left unsaid.
- **Graders that failed right answers, twice:** reworked on the kept answers before each re-run.
- **A setting no static check saw:** the skill stopped firing; the linter now warns on it.
- **Size:** 85 orders left out: lists render whole, the full text is saved as the answer of record.
- **Right figures under wrong labels:** a separate audit read the answers against the CSVs; labels are now given word for word.

What a client gets

- The owner's definitions written down once, and one recurring question pinned on their export (Stripe, Amazon or another), with a check on every answer.
- Skills for the task the team repeats: definitions asked once in plain words and printed under every answer, figures from scripts, answers checked, every excluded record named.
- The same bench on the client's skills and exports: routing on how their people ask, and their answers seeded with the defects their process can produce. Every model answer and grade is kept and handed over.

Every result and model answer these pages count, and the stores' data and truth, are public: github.com/5eBeKs/claude-skills-zoo. Later runs on stores A to C (0.9.0, 0.10.1), repeat runs of the store C comparison and earlier reviewer runs are not published; the skills, checks and bench are private. Contact: [Upwork](#). Not affiliated with Shopify, Stripe, Amazon, Intuit or Anthropic.