

Claude skills where the reader sees what every figure means

A strong model rarely gets the arithmetic wrong any more. What goes wrong is the context: which orders count as sales, whether revenue holds VAT, which month a refund or a payout belongs to. Three skills for a Shopify store's month-end ask the owner's definitions once, compute every figure with a script, check the answer against the script's results, name every order left out, and print under every answer how it was counted. A bench then asked the owner's questions on three synthetic stores, the third at a real store's size (4,470 orders over three months), with the skills and without them, and read every answer.

What the owner can see (Opus 5.5, three runs per question, 27 runs per arm)

	WITH THE SKILLS	WITHOUT
Right figures	27/27	26/27
On the skill's stated basis	27/27	13/27
Everything the owner should see, named	27/27	21/27
A "How this was counted" section under the answer	25/27	0/27
Says it added the figures up by hand	0/27	9/27

STORE	RIGHT, WITH	WITHOUT	ALL NAMED, WITH	WITHOUT
Store A	9/9	9/9	9/9	9/9
Store B	9/9	9/9	9/9	9/9
Store C	9/9	8/9	9/9	3/9

Without the skills on store C the model had a shell and wrote its own code. It read the owner's definitions file and applied it, and its figures were right in 8 of 9 runs; in the ninth its own code labelled a total as after discounts when it was before them, and August revenue came out about £1,000 high. It named the orders behind its counts in none of the 6 summary and payout runs: not the 85 orders left out of August's sales, not the 400 payout items (on the margins question it named every product it had to).

The answer zoo: it ran, and the answer is wrong

24 defects planted in correct answers (a wrong figure, the wrong rows counted, a basis named wrong, a right figure under another's name, a lost caveat, a claim no figure supports, the question not answered). The number check stopped **6**, the coverage check **5**; the other 13 went to reviewer models, three runs each.

REVIEWER	NAMED THE DEFECT	OBJECTED TO AN UNTOUCHED ANSWER
Sonnet 5	13 of 13	0 of 9
Opus 5.5	13 of 13	3 of 9

Every objection to an untouched answer was about a definition the answer left unsaid; I agreed with each and changed the answers.

The skill zoo: does the right skill run

11 defects in the summary skill's frontmatter or files, each in its own copy of the plugin, then `claude plugin validate --strict`, the linter and a routing eval (six questions, three runs each). The last column: how often the summary skill ran on its own question, asked plainly and indirectly.

DEFECT	VALIDATE	LINTER	SUMMARY SKILL RAN
none	passed	silent	6/6
vague_description	passed	W01, W02, W05	6/6
no_trigger	passed	W01, W05	6/6
overlaps_margins	passed	W04, W05	6/6
too_broad	passed	W01, W05	6/6
name_not_kebab	passed	E04	6/6
name_differs_from_folder	passed	E04	6/6
description_field_type	flagged	E03, E05	0/6
frontmatter_not_first	flagged	E01	0/6
model_invocation_disabled	passed	W06	0/6
trigger_past_the_limit	passed	E06	6/6
script_missing	passed	E07	6/6

Checking the checks

After each run every failure was read, and every pass. Each finding changed the tools:

- **A definition the skill, its truth and its graders shared:** a model without the skill reconciled to the bank; it was right.
- **A question that said less than the script did:** "test orders" meant a tag to the owner, a gateway too to the script.
- **Reviewers right about the reference answers:** every objection was a definition left unsaid.
- **Graders that failed right answers, twice:** reworked on the kept answers before each re-run.
- **A setting no static check saw:** the skill stopped firing; the linter now warns on it.
- **Size:** 85 orders left out: lists render whole, the full text is saved as the answer of record.
- **Right figures under wrong labels:** a separate audit read the answers against the CSVs; labels are now given word for word.

What a client gets

- Skills for the task the team repeats: definitions asked once in plain words and printed under every answer, figures from scripts, answers checked, every excluded record named.
- The same bench on the client's skills and exports: routing on how their people ask, and their answers seeded with the defects their process can produce. Every model answer and grade is kept and handed over.

Every result, every model answer and the stores' data and truth are public: github.com/5eBeKs/claude-skills-zoo. The skills, checks and bench are private. Contact: [Upwork](#). Not affiliated with Shopify or Anthropic.